

Care Note — Round Two

Response to the sixteen clinic scenarios and the twelve-capability list. Architecture, schema and latency are unchanged and remain in [TECHNICAL BRIEF.md](#); this covers what the review changed. Per-scenario verdicts and their tests are in [SCENARIO_COVERAGE.md](#).

Where we landed: 9 SURVIVES · 6 PARTIAL · 1 DOES NOT on the scenarios, from 6 · 6 · 4 in our own first assessment. Nineteen decisions (D-070 to D-088), 173 new tests, 528 backend test functions (851 parametrised cases) and 67 component tests passing.

Two rows moved backwards, deliberately. Scenario 3 was SURVIVES and is now PARTIAL: a patient's phone number was travelling in a query string and therefore into the access log. Scenario 2 was SURVIVES and is now PARTIAL: we had answered *where* clinic isolation is enforced and never measured what happens when that one line is wrong — the answer is every patient in both clinics, and nothing else catches it. Sections 7 and 8 cover both.

pytest tests/test_survival_scenarios.py -v walks the sixteen scenarios one at a time, and fails if its verdicts drift from the table below.

1. What the review actually found

Our own assessment came out 6/6/4 before any code changed. The pattern in those failures was more useful than any individual fix: **what survived was structural, what failed was operational**. Access control fused into a type, redaction fused into the only code that can reach a network, highlights anchored to a version — all held. What failed was every question of the form *what happens when*: when the provider is down, when the server crashes, when a clinic needs to add a patient on a Tuesday. We had built a record system and had not built the clinic around it.

Two defaults explain most of it, and both are worth stating because they generalise beyond this build.

The stub LLM provider cannot fail. It is in-process, so it cannot time out, refuse a connection or return a 503. Every test run and demo for the entire build executed against a provider physically incapable of failing — which is why a provider outage turned out to be an unhandled 500 with a traceback. The default is still right; it makes the build runnable without an API key. But it bought that determinism by removing the failure surface from view, and nothing made the trade visible. CARENOTE_LLM_FORCE_UNAVAILABLE now exercises that path on every run.

The seed script stood in for features. init_db.py runs at Phase 1 step 1, so from the first commit every test and demo began with a populated database. *How does a patient come to exist?* never arose, because patients always already did. Anything a seed provides is a feature you have not built and will not notice missing — which is exactly scenario 1.

2. Scenario verdicts

#	Scenario	Was	Now
1	Patient with no email	PARTIAL	SURVIVES — staff enrolment, phone as identifier
2	Clinic isolation, one line	SURVIVES	SURVIVES — AccessScope.query(), three test files
3	Read your logs	DOES NOT	PARTIAL — crash path fixed; a URL leak found later (§7)
4	Prove the ordering	SURVIVES	SURVIVES — one exit, asserted by source scan
5	Clinic B on Monday	PARTIAL	PARTIAL — zero schema changes; config surface built, vocabulary still global
6	Trilingual consult	PARTIAL	PARTIAL — parity for en/ms; Hokkien flagged, not read
7	Allergy at minute two	DOES NOT	DOES NOT — batch; boundary asserted by test
8	Model hangs 45s	PARTIAL	PARTIAL — 8s timeout and cancel; no server-side abort
9	Provider 503 for an hour	DOES NOT	SURVIVES — degrades to extractive, and the card says so
10	Two people, same note	SURVIVES	SURVIVES — version check plus unique constraint
11	Link never received	DOES NOT	PARTIAL — reach modelled; no sender exists
12	Summary wrong by one dosage	PARTIAL	SURVIVES — correction banner and a dose gate
13	Allergy asserted vs denied	PARTIAL	SURVIVES — assertion_vs_denial at HIGH, one card per disagreement
14	A number that means nothing	SURVIVES	SURVIVES — floor now language-independent
15	Ranking learns from what it showed	SURVIVES	PARTIAL — floors, exploration and inspectability hold; bias now measured at 0.71 rather than argued (D-092)
16	Highlight cites edited source	SURVIVES	SURVIVES — now side by side, both versions named

On the twelve capabilities: 4 SURVIVES · 7 PARTIAL · 1 DOES NOT, revised down from 6 · 5 · 1 by a second audit (§10). The single DOES NOT is streaming ASR. Full per-row detail with the tests that cover each is in [SCENARIO_COVERAGE.md](#).

Four rows moved down, and the reasons differ. Two were overclaims: the correction leg of fact extraction was justified by “version history as the correction record”, which is a human editing a note and not a speaker correcting themselves in a transcript; and exposure bias was marked SURVIVES on the strength of a mitigation existing, when the capability asks for an evaluation. One hid a defect (§10). The fourth, audience-appropriate output, is a capability we declined on purpose — no machine-written text can become patient-facing at all (D-067) — and ticking it would have claimed a feature where we had made an argument. Stated as a refusal it is the stronger answer; stated as a tick it is the thing the feedback said counts against us.

3. What changed

Failure handling (3, 8, 9). The chokepoint now translates timeouts, transport errors, 5xx and 429 into one LLMUnavailableError; it does not decide what to do about them. Each caller chooses, because the right answer differs by purpose: the scribe degrades to the deterministic extractive summariser it already had — labelled offline-extractive-v1:provider-unavailable, so a degraded note is legible as degraded — and a patient-facing generator would refuse. 4xx stays loud: a bad API key hiding behind a slightly worse summary is indistinguishable from an outage. Timeout 60s → 8s, plus a cancel button, since the transcript is stored before the model is called and abandoning loses the summary, never the consult.

Crash logging is a separate chokepoint: type, route and an eight-character reference, never str(exc), which is where SQLAlchemy puts bound parameters.

Language parity (6, 14). The risk floor now works over canonical tags rather than English strings, so it inherits every language the tagger knows. It also gained negation handling, asymmetrically: a symptom is dropped only if it appears exclusively inside a negation, so “no chest pain Monday, chest pain today” still rates high, and a denied symptom drops to medium rather than to nothing. Content in an unsupported language is flagged as unread rather than silently producing nothing — abstention beats confident silence.

Denials as claims (13). Negated allergy mentions become allergy_denial claims instead of being discarded, and assertion_vs_denial fires at HIGH — not CRITICAL, because nothing dangerous has happened yet and diluting CRITICAL would break the level that means *someone is about to be given a drug they react to*.

Reach (11, 12). unread / read / corrected, derived from timestamps PatientView had recorded since D-033 and nothing had queried. dispatched is deliberately absent because nothing dispatches. The patient view leads with a plain-language correction banner, computed before the read marker moves.

Enrolment (1, 5). Staff can register a patient and issue a login; identifier type is explicit and a phone number is first-class. Patient.dob and mrn became nullable — both were NOT NULL, which was a second and quieter way the schema decided that someone known only by a phone number was not a patient.

Regeneration (capability list). Was undefined behaviour, found by probing: a new session id produced a duplicate summary, the same one crashed on a unique constraint. Now it reuses the entry and appends a version, so accepted highlights and comments survive — and refuses outright when a human has edited, because merging would mean choosing which of a clinician's sentences to keep.

Dosage (capability list, and the hint). Doses were compared against each other and never against a reference, so the build could see two entries disagree and not see 5000mg. A 17-drug adult reference now bands doses plausible / unusual / implausible; only implausible gates, and only on patient-facing writes, with the override recorded. Acknowledgement rather than refusal: a hard block teaches people to route around the check.

Provenance (16). Stale highlights now show the anchored text and the current text side by side, with both version numbers named.

4. What we tried that did not work

An order-of-magnitude dose threshold. Metformin 5000mg passed as merely unusual — the exact case the hint describes. The fix is principled rather than tuned: ranges are single-dose, a legitimate daily total reaches ~3× that (TDS), so beyond 3× exceeds any plausible daily total. 1500mg BD correctly stays unusual.

@app.exception_handler(Exception) for the log leak. Does not work. Starlette's ServerErrorMiddleware calls the handler and then re-raises so the server can log the traceback — sanitising the response and leaving the leak untouched. It had to be middleware. We only found this because the test asserted on captured log output rather than on the

response body.

A blanket-denial pattern that matched too much. `denies allergy to aspirin` registered as a denial of *all* allergies. A contradiction detector that cries wolf is worse than one with gaps: the gap loses a finding, the false positive teaches people to stop reading all of them.

Streaming capture — not attempted. Retrofitting incremental extraction is tractable; `test_capture_timing.py` asserts the deterministic extractors work on two turns alone. What we could not resolve is the interruption policy. When it is acceptable to interrupt a doctor mid-consultation is a clinical workflow decision, and guessing would have produced a demo rather than a product.

5. Assumptions that no longer stand

"Logging is solved because `log_event` is content-free." Solved for content we log on purpose; silent about content logged on our behalf. One unhandled 500 put a name, an NRIC and note content in a single line. The earlier brief claimed "logs grepped, zero hits" — true of the success path, and the wrong question.

"Multilingual is fixed because the tagger handles Malay" (D-058). The fix landed at the tagger; the risk floor was a separate path making the same English-only assumption. Every Malay fixture passed because they only exercised the tagger. *When you fix an assumption at one layer, go looking for other layers that made the same one.*

"Dropping negated mentions is correct." Correct for its purpose, and it discarded the patient's denial entirely — a denial is a position, not an absence.

"The medication cue vocabulary is adequate." Built against written notes. *I will start you on amoxicillin* — how prescribing is actually said — classified as a bare dose claim and never reached the allergy comparison.

Still standing: regex redaction over NER for auditability (D-012); no HTML-escaping on write, because corrupting dose `<5mg` is worse than the XSS it prevents (D-015); full snapshots over diffs (D-006); least-privilege on staff visibility (D-004); and reporting contradictions without ever resolving them (D-068) — tested hardest, and the one we would defend most strongly.

6. Where we expect this to fail next

1. **A consult beyond English and Malay.** The unread flag means the system says so, but the content is still invisible downstream. First week.
2. **A dose the 17-drug table does not know.** No check at all, and the absence looks identical to a pass.
3. **Alert fatigue on contradictions.** Four classes now fire; nothing monitors the dismissal rate, which is the signature NEVER_DAMPENED exists to survive.
4. **Regeneration refusing too often.** Any human edit blocks it, including a typo fix — possibly annoying enough to teach people not to edit AI summaries.
5. **Two disagreeing clinicians in the learning loop.** Untested rather than designed; saturation bounds the damage.

Still not built, plainly: streaming ASR, acoustic diarization, any message sender, per-clinic configuration, presence indicators, and a real drug database. Scenario 7's *timing* remains DOES NOT — nothing is incremental, and `test_capture_timing.py` asserts that boundary rather than papering over it. Its *detection* half was a defect rather than a boundary, and is fixed (\$10).

7. Re-auditing our own answers

Everything above was written, then the build was probed again as if by someone trying to disprove it. Three defects came out, **all inside rows already marked SURVIVES**, and none caught by the 509 tests passing at the time.

The common cause is worth more than the individual fixes: **each test used the minimal shape of the case it was testing**, and the minimal shape is the one where the bug cannot appear.

Contradictions fanned out N×M (D-081). Every contradiction test used one assertion and one denial. That is the single shape where pairwise output and grouped output are identical. A real chart re-records an allergy at every visit, so four "allergic to penicillin" entries against two "no known allergies" entries produce eight findings that all say one thing. The Glance View caps the list at five, so the copies filled the cap and an unrelated **metformin dose disagreement was evicted from the card entirely** — a real unresolved conflict made invisible by a different one being mentioned more often, degrading as the record grew. Detection stays pairwise; the display unit is now (kind, subject) and every supporting entry keeps its own pointer, because a card reading "and 7 others" without links would trade an alert-fatigue problem for a provenance one.

Degradation was legible to an auditor, not a clinician (D-082). Section 3 claimed a degraded note is "legible as degraded." True of the database. In the interface the only trace was `ai_model_used` rendered as a 10px grey monospace string in the provenance footer, next to the pointer — a machine-facing identifier in the place on the card a clinician looks least. During an hour-long outage the card read as an ordinary AI summary. `ai_degraded` is now a wire field with a chip in the existing vocabulary, kept deliberately separate from the confidence signal: "the model was unsure" and "no model read this consult" call for different responses.

A phone number in a query string (D-083). `POST /patients/{id}/login` took the patient's identifier as a query parameter. For scenario 1 that identifier *is* a phone number. The ASGI access log records the full request line — path and query — before the application sees the request, and identically for requests that 404 or are refused by RBAC, so neither `log_event` nor the sanitised error middleware touches it. **The feature built to answer scenario 1 leaked through the door scenario 3 asks about.**

That last one is why scenario 3 is now PARTIAL. Our answer had been assessed against the crash path — the door we had just finished guarding — and not against the sink that logs every request whether or not our code runs. The parameter moved into the request body, and `tests/test_url_surface.py` now asserts across every route that path and query parameters come from an explicit allowlist of opaque ids, enums and structural pointers, with a second test asserting the allowlist itself stays free of PHI-shaped names so the invariant cannot be satisfied by quietly widening it.

What this says about the rest of the table. Three overclaims in rows we had tested and believed, found in one pass. We would expect a fourth to exist. The rows we would probe next are 12 and 15 — the dosage gate has only ever been exercised against drugs the 17-entry table knows, and the learning loop's NEVER_DAMPENED floors have never been tested against a clinic that dismisses at a realistic rate over months rather than in a simulated burst.

8. The themes underneath the sixteen

Read as eight themes rather than sixteen incidents, the scenarios sort our build into two piles, and the split is not where we expected.

Enforcement that is strong but singular (theme 2). Access control is fused into a type and redaction into the only module that can reach a network. Both are impossible to forget. Neither is defended in depth: dropping the clinic predicate from `AccessScope.query` exposes every patient in both clinics and nothing catches it (D-085). We had been treating *unforgeable* as if it were *redundant*. It is a different property, and the scenario asks about the second. The database-level predicate check that would give real depth is scoped in D-085 and deliberately not attempted this close to the deadline.

Privacy outside the front door (theme 3). Redaction-before-LLM is provable (`test_llm_chokepoint.py` — no module but the wrapper may reach a model). The windows were the problem: crash logs carrying SQLAlchemy bound parameters (D-071), then the access log carrying a phone number we put in a URL ourselves (D-083). Both were found *after* we had declared the door guarded. The pattern worth naming is that each leak was in a sink our own code does not write to.

Numbers that mean something (theme 6). Risk has a deterministic floor a model cannot lower, and the floor is language-independent (D-072). Confidence is derived, not self-reported. Extraction and generation are separated in the code, not just the prose: highlights are character offsets into a pinned `source_version_number`, never model-paraphrased text, which is what makes scenario 16 answerable at all.

Feedback loops (theme 8). This is where re-reading the themes changed the build most. NEVER_DAMPENED floors a protected tag's learned weight, and we had been citing it as the answer to "what stops the ranking burying an allergy." It floors the wrong quantity: surfacing is a top-N cut, so other tags rising displaces an allergy with its own weight untouched, and one dismissal removed it from the card permanently. Protected classes now bypass ranking entirely (D-084). Learning orders the protected set; it no longer decides membership.

Identity and tenancy (theme 1), answered specifically. Scenario 5 asks for config-versus-schema and the answer is now precise. *Schema: nothing* — every table carries `clinic_id`, isolation is enforced and tested both directions, and the seed has run two clinics since Phase 1, so a third needs no migration. *Config: everything, which was the problem* — every value a clinic might differ on was a module constant, so "Clinic B keeps records for a year" was a deploy. `ClinicConfig` (D-086) makes volume and retention per-clinic, with an absent row meaning defaults so onboarding stays zero-step.

The more useful half was deciding what stays global. Redaction patterns, the protected highlight classes from D-084, contradiction severities and the dosage reference are all asserted *not* to be configurable, under one rule: **a clinic may change what it sees, never what it is protected from**. Anything that could be turned down until an alert stops firing sits on the left of that line. The gap that keeps scenario 5 at PARTIAL is clinical vocabulary, still global — and it is the piece needing most care, because additions must be additive only or a clinic could remove `entity:allergy` and reach the safety floor sideways.

Where we are still weak, plainly. Degraded-mode behaviour (theme 4) is handled for the model and absent for delivery — there is no sender, so scenario 11 cannot fully survive. Concurrency (theme 5) loses no updates but is not real-time. And scenario 7 remains DOES NOT: the scribe is post-hoc by construction, so a drug allergy at minute two is not in the Glance View until the consult ends. That is an architecture decision, stated rather than hedged, and `test_capture_timing.py` pins the boundary rather than papering over it.

9. Two failures the clinician sees, and the sinks we had not scanned

A last pass over ease-of-use and leakage, asking only “what does a user experience when this goes wrong” rather than “does the logic hold”.

A render crash white-screened the app. There was no error boundary anywhere, so any component throwing unmounted the whole tree. A clinician with a patient in the room got a blank page — and in a clinical record a blank page is indistinguishable from data loss, so the recovery they reach for is retyping a note they never lost. The boundary now leads with *nothing you saved has been lost*, which is true because writes commit server-side before the response returns. It renders no error text at all: `error.message` can carry interpolated note content, and a stack trace on a shared consult-room laptop leaks in front of the patient.

A session timeout stranded the user. Sessions last 60 minutes with no refresh flow — a deliberate security decision (D-016) that fires on a real clinic laptop most afternoons. The 401 surfaced as a red line reading “Token expired” beside a chart that had silently stopped working. A security control that strands people is a control people route around, so it is now handled centrally: the sign-in screen returns, states that saved work survived, and the session-restore probe is exempted so a first-time visitor is not told their session expired.

The sinks we had not scanned. Server logs were governed by `log_event`, crash output by D-071, URLs by D-083 — and the browser console by nothing. One `console.log(entry)` left in during debugging puts a full note somewhere most deployments forward to a third-party dashboard, and the app looks identical either way. A scan now fails the build on any console statement outside two content-free allowances, and the error boundary — the one file that logs on a failure path — is itself pinned to logging a type name rather than a message.

What this pass confirmed rather than fixed. Patient-role isolation holds across every surface we could reach: nine endpoints probed with a patient token, zero fragments of staff notes, clinician sections or raw AI summaries in any response body. The service worker is network-only for `/api`, so no patient data is written to disk on a shared device. Neither needed changing, and both are worth stating because “we checked and it held” is a different claim from “we assumed it held”.

Known gap. React re-logs caught errors to the console in development builds, message included, and a boundary cannot suppress it. `Resilience.test.jsx` asserts that the leak happens, so it is visible in the suite rather than only in a document, and it will fail if React changes the behaviour.

And the one that is ordinary rather than exceptional. This is a PWA for bedside use, so losing the network is not an edge case. Every such failure reached the clinician as the browser’s own words — “Failed to fetch” — which is Chrome-specific, unstable across engines, and silent on the only question that matters mid-consult: did the note save. Reads and writers now say different things, because the reassurance differs: a failed read changed nothing, and a failed write did not save but the text is still in the box. That second claim is true rather than hopeful — draft state already survived a failed save, which this pass confirmed rather than built. An offline blip also deliberately does not sign anyone out, so a nurse who walks into a lift does not come out of it retyping a note she already wrote.

What is still missing there is a queue. Nothing retries, nothing is held, and the honest reason it was not attempted is that a durable outbox means unencrypted patient text sitting in browser storage on a shared ward device — a data-protection decision, not an afternoon of work.

10. A second audit, and the pattern behind all six defects

Section 7 reported three defects found by probing rather than re-reading. We ran the exercise once more against the twelve-capability list specifically. Three more came out — and the interesting result is not the individual bugs but that **all six share one cause**, which we can now state precisely.

Contradiction detection could not see inside one entry (D-089). `detect()` compared entries pairwise and never an entry with itself. For typed notes that is correct and deliberate. But `run_scribe` writes **one Entry per consult**, so a twenty-minute conversation is one row — and an allergy at minute two against a prescription at minute nineteen was undetectable at any point in its life. Not during the consult (nothing is incremental, which we had documented) and not after it either, which we had not noticed. `test_capture_timing.py` passed the whole time; it asserted the timing boundary and was blind to the detection gap sitting beside it. Intra-entry comparison is now enabled, gated to three classes, with dose corrections requiring an explicit retraction cue so that clinical deliberation (“500 or 1000, depending on tolerance”) stays quiet.

The abstention flag never fired for unspaced scripts (D-090). `is_unreadable` is the whole of our answer to “the transcript is trilingual”: when the tagger cannot read a turn it says so out loud, rather than producing an empty tag list that a clinician cannot distinguish from “nothing clinical was said”. It measured substantiveness as `len(text.split()) >= 6`. Chinese and Japanese are written without spaces, so a whole paragraph is one token, falls under the bar, and returns `False`. The comment beside our Malay vocabulary names Mandarin as a known-uncovered language and rests on the abstention flag catching it. It did not. **A documented gap that a second mechanism is silently failing to cover is worse than an undocumented one, because the documentation reads like a control.**

Exposure bias had a mitigation and no measurement (D-092). Now measured: displacement 0.29, exposure concentration **0.71**, blind-tag rate 0.31, zero protected classes displaced. The 0.71 is the bias stated as a number — five of seven visible slots go to tags this clinic has already given feedback on. We are not claiming that is a good number; we have nothing to compare it to. We are claiming it is a number that moves when the system changes, which is what the capability asks for and what an argument cannot do.

We fell into the same trap while building the measurement. The first version of the evaluator ranked by score and took the top N. That is not what our Glance View does — D-084 surfaces protected classes regardless of rank — and measured that way the report claimed entity:`allergy` never reaches the card. `False`, and alarming enough that it would have gone in this brief as a finding. It is caught by a test that fails against the naive implementation.

One finding we did not fix, because the fix is bigger than the bug. Keyword tagging has no notion of negation — “no anaphylaxis” emits `symptom:anaphylaxis` — which is a known limitation, pinned by a test, with a stated reason (changing it moves English scoring and needs the Glance View re-measured). What was *not* documented is its interaction with the protection list. `symptom:anaphylaxis` is in `NEVER_DAMPENED`, so a note explicitly ruling anaphylaxis out produces a tag that can never be learned down: the clinician dismisses it, and the floor discards the dismissal. The anti-alert-fatigue mechanism manufactures a permanent false positive. The safe-direction argument we made for negation (“a ruled-out symptom is surfaced for a human to dismiss”) holds for ranking and does not hold here, because on this path the dismissal has nowhere to go.

The pattern, stated once. Every one of these six defects was invisible to its own tests, and in each case for the same reason: **the test used the shape of the case its author had in mind**. Cross-entry contradiction tests used two entries. Abstention tests used romanised Latin script. Exploration tests passed `existing=[]`. The evaluator modelled the card its author assumed. None were careless; each was written from inside the assumption it needed to escape, which is not a thing more of the same tests would have fixed.

So we changed their shape rather than their number, and it found two more. Property-based generation over redaction input, and enumeration of the clinic boundary from the live OpenAPI schema rather than from a hand-picked list.

A phone number could hide behind an en-dash (D-095). Every separator class in `redaction.py` is spelled in ASCII. `\s` is Unicode-aware so exotic spaces cost nothing, but the dash class is not — `hp 9123–4567` passed through untouched, and `find_residual_phi` reported clean. The second half is the real finding: that function is both the fail-closed tripwire and the oracle our property tests assert against, and it shares its regexes with the redactor. **A check and its own test must not share an implementation**, or the gap is invisible twice. Both now fold separators. iOS autocorrects hyphens between digits into en-dashes, so for a build whose premise is text arriving from phones, this was the ordinary path rather than the exotic one.

Clinic isolation is now enumerated, not sampled (D-096). Against the exact mutation the feedback describes — deleting the `clinic_id` filter — the hand-written RBAC tests raise 15 failures and the enumerated matrix raises 48. Both catch it; only one tells you what it exposed. Building it also showed that sending an empty body made eleven write routes return 422, because FastAPI validates before the RBAC dependency runs — a test counting that as “refused” would have passed for the wrong reason on every write route we have.

Four more properties held, and that is also a result (D-097). Content round-trip through the real API: `BP <120/80` and `<script>alert(1)</script>` both come back byte-identical, so neither promise in D-015 is being satisfied at the other’s expense. Analyser totality and determinism over the full unicode range — an unstable score cannot be wrong, because it never says the same thing twice. Revision-history invariants over generated `edit/revert` sequences, of which the load-bearing one is that already-written versions are frozen: highlights anchor to `source_version_number`, so a mutable version would make every provenance pointer a liar. And log hygiene, greping every record from every logger at `DEBUG` — including exception text — for synthetic identifiers written through `create`, `edit`, `comment`, `refusal`, `validation failure` and the `scribe pipeline`. Nothing leaked.

Every one was mutation-checked, because a test that has never failed is a test whose teeth are unmeasured. Deleting the `clinic_id` filter, rolling the version number backwards on `revert`, and adding a single `logging.info(“updating: %s”, payload.content)` are each caught by the property that claims to cover them.

We also found a leak we chose **not** to fix: space-separated identifiers `900101 01 5432`, which is what a patient reading an IC aloud transcribes to) survive redaction. Widening the pattern puts every run of grouped clinical digits at risk of becoming `[ID_1]` — trading a narrow privacy gap for a broad accuracy one, which is the wrong direction and exactly what the hint warned about. Pinned by a test that asserts current behaviour and fails the day anyone changes it.