

Care Note

Nightingale 72-Hour Build. Synthetic data only; not safe for real PHI as-is (§7).

1. What Care Note is, and how we approached it

A patient's history is currently spread across dated free-text notes, each written by a different person in a different screen, and there is no single place that consolidates all of it. This makes opening a chart before a consult difficult, because the clinician has to reconstruct months of history by scrolling through notes and working out for themselves which ones matter.

Care Note replaces that with one shared record per patient. Clinician notes, nurse and staff notes, the patient's own contributions, and AI-scribed summaries of every consult all go into a single timeline, and each entry records who wrote it, when, and what it came from. Above the timeline is the **Glance View**, a card designed to be read in under ten seconds. It answers four questions: what changed since you last looked, what could hurt this patient, what matters and why, and what is outstanding and whose job it is.

Four things had to be true for this to be worth using, and we designed around all four rather than picking one.

It has to be secure. These are medical health records. Access control checks both the user's role and their clinic on every request, and it runs on the server, so the UI is never what protects the data. Any text sent to a model first passes through one redaction function that removes names, ID numbers and phone numbers. That function then re-checks its own output and refuses to send anything that still looks identifying. Logs record IDs, actions and timestamps — never the content of a note.


It has to be fast. Clinical staff see a lot of patients in a day, so a tool that adds thirty seconds per chart costs them hours a week. We measured this rather than assuming it: the Glance View loads with a P95 of about 11 ms against the 300 ms target (§4).

It has to be genuinely useful, or nobody will adopt it. Software that makes a clinician's job harder does not get used, whatever it does well. So the design removes work instead of adding steps. The AI writes the consult summary that somebody would otherwise have typed. Confirming or dismissing a suggestion takes one click and does not navigate away from the card. Each role sees the part of the record that concerns them rather than all of it. The patient view is written in plain language rather than clinical shorthand, because the person reading it is often anxious and is not a clinician.

And the AI has to be verifiable, continuously. This is the constraint that drove most of the architecture. Neither kind of error is acceptable here: a false positive teaches staff to ignore the system, and a false negative can mean a missed allergy. We do not think a better model fixes that on its own, so the system treats **AI output as a claim rather than a fact**. Every claim has a source the clinician can open, a confidence score measured from the transcript rather than reported by the model, and a human who can accept or reject it in one click — and it can never overwrite something a clinician wrote. When the system cannot back a claim up, it says so instead of guessing: a summary line it cannot trace to the transcript is shown with no citation at all, rather than one that looks checkable and is not. This is a decision about the data model before it is one about the interface, which is why it lives in the schema (§3).

2. Architecture

A React frontend, a FastAPI backend and a SQLite database. Two parts of the diagram matter more than the rest. `require_access()` is the only way a request reaches a handler, so every access-control check happens there. Only `llm_client.complete()` can send text to a model, so every redaction check happens there. Both are deliberately single points: one rule that nothing can bypass is safer than the same rule repeated correctly across forty handlers.



Python 3.11 / FastAPI / SQLAlchemy / SQLite (Postgres-ready), React 18 + Vite + Tailwind, pytest. FastAPI specifically because dependency injection is what makes the RBAC rule unforgettable:

RBAC fuses role and clinic inseparably. The common real failure is a route that checks role and forgets the tenant filter. Here that is not expressible. `require_access()` is the only way a route learns its caller — there is no exported `get_current_user`. It yields an `AccessScope`, never a `User`, and that is also the only DB handle a route gets. `scope.query()` applies the clinic predicate itself and **raises `TypeError`** on any model lacking `clinic_id` rather than returning it unfiltered. `clinic_id` comes from the verified JWT only. No handler in `patient_routes.py` mentions `clinic_id` in a filter. Cross-clinic fetches by exact id return **404, not 403** — a 403 confirms the id exists.

One redaction chokepoint. `redact_phi()` is called unconditionally by `complete()`, the only module that reaches a model. Callers cannot opt out; redaction is idempotent. After redacting, the payload is re-scanned and the call **raises rather than sends** if PHI survived. Three source-scanning tests fail the build if any other module imports an LLM SDK, references a model endpoint, or removes the redaction call.

3. Schema

Entry is one item on the patient's timeline, and it is the centre of the schema — every other table connects to it. Versions hold the revision history, Comments hold the collaboration, Highlights are what the Glance View displays, and an AIScribedNote row is what marks an entry as written by the AI rather than a person. Provenance works differently from the rest on purpose: it is a string URI instead of a foreign key, because a citation can point at a whole entry, a range of characters inside one, or a segment of audio, and those are not all rows in the same table. How the learning mechanism fits in is covered at the end of this section.

Link	Mechanism
Entry → Versions	<code>Version.entry_id</code> , unique <code>(entry_id, version_number)</code> . Full snapshots, not diffs — revert is a copy, and no chain can be corrupted by one bad link
Entry → Comments	<code>Comment.entry_id</code> ; threads via self-referential <code>parent_comment_id</code> ; open/resolved status
Entry → Highlights	<code>entry_id</code> + character span + the source_version_number the span was computed against , so an edit cannot silently move a highlight onto different text
Entry → AIScribedNote	One-to-one. This row's <i>presence</i> is what makes an entry AI-authored; <code>author_role='system'</code> is the denormalised fast check. A UI cannot render an AI note as a clinician's
AIScribedNote → transcript	Shared <code>session_id</code> with <code>TranscriptSegment</code>
Summary line → spoken words	<code>SummaryAttribution</code> links each summary line to the segment that produced it
Anything → Provenance	A string URI , not a foreign key

CLINIC <- USER, PATIENT, ENTRY, FEATURE_WEIGHT · PATIENT <- ENTRY, TASK, CAPTURE_SESSION · ENTRY <- VERSION, COMMENT, HIGHLIGHT, TASK, SUMMARY_ATTRIBUTION and -o AI_SCRIBED_NOTE, ENTRY_ARCHIVE · COMMENT <- COMMENT, TASK · USER <- INTERACTION_LOG, AUDIT_LOG, PATIENT_VIEW. Full Mermaid ER diagram in `SCHEMA.md`.

Provenance is a URI — `entry://<id>#span:<start>-<end>`, `session://<id>#turn:<n>`, `transcript://<id>#segment:<n>`. A foreign key points at one table; provenance targets are heterogeneous — a whole entry, a character range in one, a turn in an AI session, an audio segment that is not a row here. One resolvable string covers all and keeps "click a highlight, land on the source" to one code path. Cost: the DB won't enforce integrity, so `resolve()` is the only dereference path, it **raises on a dangling pointer** rather than degrading to empty, and it enforces `clinic_id` — a valid pointer must never become a cross-tenant read primitive.

clinic_id is denormalised onto every scoped table, even where a join would derive it. That redundancy is what lets `AccessScope.query()` apply one uniform predicate to any model — and why it can *refuse* one lacking the column.

Learning integration. `InteractionLog` records what clinicians touch as **extracted tags only, never prose**. `record_interaction()` calls `apply_signal()` in the same operation, so no route can log a signal the learning table never sees — the same chokepoint reasoning as redaction. Evidence aggregates per `(clinic_id, feature_tag)` into

FeatureWeight with a 90-day half-life, saturating into (-1, 1), read by scoring.Learned_component() as one term. FeatureWeight is a materialised view of the log, never nudged, so it cannot drift from its evidence. It is capped at 0.25 of the total, cannot invent a highlight (the rule layer runs first), cannot cross a clinic, cannot be trained by patients, and **cannot silence safety vocabulary** — allergy, sepsis, anaphylaxis and self-harm tags are floored at zero.

4. Glance View and latency

The Glance View is the main thing we are claiming: that a clinician can get oriented on a patient in under ten seconds. Two decisions make that work, and both involve doing less. The card deliberately shows only a few items rather than everything it could — a card showing everything would just be the timeline again — and the expensive scoring happens when an entry is written, not when the card is opened, so loading a chart is mostly reading rows that already exist. The timing below is evidence for that second decision rather than a claim about raw speed.

The Top Card answers four questions in fixed order: what changed since you were last here; what could hurt this patient; what matters and why; what is outstanding and whose it is. Ranking is a weighted sum over named features, each highlight showing its own arithmetic and a one-line risk_reason. **Nothing is ranked by a model.**

Measured: P95 ≤ 300 ms target. A middleware reports X-Response-Time-Ms per request — request in, queries run, payload serialised, response out. 200 iterations after 20 discarded warm-ups, 11-entry chart with 6 highlights, SQLite on local disk. Re-measured after the Phase 7 fixes touched every timestamp in this payload; three consecutive runs: **P95 11.54 / 11.09 / 11.63 ms**. Range reported, not the best run. These are lower than the 13.30–15.94 ms recorded before those fixes, which is container load rather than an improvement anyone engineered — attributing it to the change would be reading noise as a result.

This **excludes network transit and browser render** — those depend on deployment and device, and folding loopback in would invent precision. The client measures its own round trip and the header shows both, so the demo never conflates them.

The figure is evidence for something narrower than "fast": **application work is a small fraction of budget** and no N+1 hides in the hot path. Two decisions carry that — **highlight scores are computed on write, not read**, so the view reads precomputed rows; and **timeline enrichment is batched** into four grouped queries regardless of chart size. Production means Postgres over a network, hundreds of entries and concurrent load; the ~20× headroom makes inversion unlikely, but the test that settles it is a loaded staging environment.

5. Trust calibration, evaluation and abstention

This section answers the question a clinician would actually ask: why should I believe any of this? Below are five mechanisms, each addressing a different way an AI-assisted record loses people's trust — the AI presenting a guess as settled, uncertainty that is not visible, human work being quietly overwritten, two notes contradicting each other with nobody noticing, and generated text being shown to a patient. The table after them takes the three numbers we put on screen — a risk badge, a confidence label and an importance score — and asks not what they are but how anyone would know if one of them were wrong, and what the system does when it cannot answer.

1. Accept / reject. Highlight.status starts suggested and needs a clinician decision; no AI claim reaches the card as fact on its own authority. One click, inline, no navigation, immediate confirmation — because this decision is also the training signal, and a high-friction control starves the loop. The interaction cost is a design constraint, not a nicety. A clinician's own hand-marked span is recorded accepted on creation and carries a scoring bonus, so human judgement outranks machine suggestion on the card by construction — which is precisely what silently broke in Phase 7 (§6).

2. Visible confidence, derived not asserted. Confidence is computed from hedging density in the source transcript on **both** paths — a live model's self-report is stored for calibration and is never what the clinician sees, because a model's opinion of its own reliability is not evidence about it. Bands are numeric and defined once: **high ≥0.75, medium 0.60–0.75, low <0.60**, and the chip shows the word and the number together. Low-confidence summaries flag **separately from risk** — "this might be dangerous" and "this might be wrong" are different warnings a clinician acts on differently.

3. Clinician precedence and a review flag. The brief allows either; we do both. A clinician edit wins immediately — care is never blocked on a resolution workflow — but conflict_flagged is set and supersedes entry_id records what was overridden, so the disagreement stays visible. Precedence alone loses information: that the AI disagreed is itself clinically interesting, and discarding it quietly is how a system teaches users to stop trusting it. AI notes are never edited in place; corrections supersede.

4. Contradiction detection, including human-human. Mechanism 3 handles a disagreement that has been resolved. The dangerous one is unresolved: a nurse records a penicillin allergy, a clinician prescribes amoxicillin, neither is wrong on purpose, and no precedence rule applies because both are people. services/contradictions.py detects three classes deterministically — allergy against administration (including by drug class), dose disagreement on the same drug, and started-here/stopped-there — cites **both** entries, and **resolves nothing**. Deciding that the more recent note wins would silently discard an allergy recorded last year. It sits above everything else on the card.

5. Patient-facing generation is structurally impossible, not approved. Showing a clinician a hallucinated line is a bad day; showing a patient one is a different category of harm — no second reader, no provenance rail, no basis to doubt. Rather than generating patient text and requiring sign-off — a step people click through under load — no generated text can become patient-facing at all. PATIENT_FACING_TYPES is writable only by clinician; Role.SYSTEM can write nothing; assert_never_patient_facing() runs at import against the scribe's own type map so a future edit fails loudly; and AI types are absent from the patient's viewable set. A clinician may read an AI summary and write an instruction from it — they type the words.

Supporting all five: provenance_pointer is non-nullable on Highlight, resolution lands on the **character span** not the note, and AI-vs-human is carried by four independent signals (rail style, colour, typeface, label) so it survives in greyscale.

What each number means, and how we would know it was wrong

	Risk badge	Confidence label	Importance score
What it is	Ordinal none → critical. Deterministic rules set a floor; a model may raise it, never lower it. The floor works over canonical tags, not English strings, so it inherits every language the tagger knows (D-072)	A number in 0.35–0.90 measured from hedging density in the source, banded high/medium/low	Weighted sum of named terms (recency, risk, entities, open actions) plus a learned term capped at 0.25
How we'd know it was wrong	Same input must give same level every run; a floor test asserts chest_pain cannot resolve below high, and a paired test asserts Malay and English forms of the same symptom produce the same floor. Drift is visible because model_proposed_risk and risk_floor_applied are both stored	Confidence must fall as hedging rises — asserted against a plain and a hedged transcript. The band boundary and the UI's flag threshold are asserted to be the same constant	Every highlight shows its own arithmetic; a wrong ranking is inspectable term by term rather than argued about
What happens when it is	The rule floor holds the badge up. risk_floor_applied tells the clinician the level came from a rule, not a model's mood	Below 0.60 the summary is flagged "verify against source" separately from risk, and the provenance rail opens the transcript	The clinician accepts or rejects in one click; rejection dampens the tag — except safety vocabulary, floored at zero so fatigue cannot silence anaphylaxis
When it abstains	Unparseable model risk falls back to low, with the deterministic floor still applied on top. A symptom that appears only inside a negation does not set the floor — but one un-negated mention anywhere does, so "no chest pain Monday, chest pain today" still rates high (D-072)	Never claims 1.0 — a summariser reading a transcript it did not hear has no business reporting certainty	A span with no clinical reason produces no highlight at all ; the rule layer runs before scoring

The same discipline governs the two places abstention matters most. **Attribution** classifies each summary line verbatim / derived / **nothing** — a line the model composed or invented gets no pointer and the UI says "no traceable source" rather than pointing somewhere plausible, because a false citation that looks checkable survives review. **Redaction** re-scans its own output and **raises rather than sends** if anything identifying survived; it is also tested for the opposite failure, that Metformin 500mg BD and HbA1c 8.2% survive intact, since over-redaction corrupts the note it was protecting.

Exposure bias is the one hazard we can only partly answer. Learning sees feedback only on spans it chose to surface, so a tag below the cut is never shown, never accepted, never weighted. One suggestion slot per entry is reserved for a candidate carrying a tag the clinic has never given feedback on — deterministic rather than epsilon-greedy, because a card that differs between loads is worse on a clinical surface than the bias it fixes. It narrows the loop; it does not close it.

6. Trade-offs, assumptions, deferred scope

Each decision below had a reasonable alternative we did not take, so we have given the reasoning rather than just listing what was built. In a 72-hour project what was left out, and why, says more than the feature list does. DECISIONS.md in the repo has the full log — around seventy entries, including scope we deliberately cut and a few decisions we later reversed. The section ends with the defects we found after we thought the build was finished.

Regex redaction over NER (D-012). Data is synthetic, so recall against real name diversity isn't what's tested — the boundary's un-bypassability is. Regex is auditable line by line, worth more in a trust system than F1 from an uninspectable model. Production layers NER behind the same signature.

Content is deliberately not HTML-escaped on write (D-015). Clinical prose contains BP <120/80 and dose <5mg. Escaping on write double-escapes on render; tag-stripping can eat <5mg and silently turn a dose limit into mg. Corrupting a note is worse than the XSS it prevents — because untrusted content is never rendered as HTML at all, enforced by a build-failing source scan.

Assumptions where the requirements were silent: staff cannot view `clinician_sections` (D-004, least privilege); admin reads all in-clinic but authors no clinical content (D-011), so it cannot quietly alter the record.

Optimistic locking over CRDTs. RBAC already partitions who writes what, so most conflicts are prevented by construction; presence is polish paid for in infrastructure.

Multilingual: partially closed, and the design point matters (D-058). Phase 5 carried code-switched speech through intact but tagged English only, so identical clinical content produced `['symptom:swelling']` in English and `[]` in Malay — no tags, no score, never on the Glance View. That fails in the worst direction: the patients least likely to be understood in English are the ones the system stops surfacing. A Malay vocabulary now maps each term to the **canonical English tag**, so `bengkak` emits `symptom:swelling`. That is not cosmetic — tags are the keys Phase 4 learns against, and a separate `symptom:bengkak` would have split one concept into two features and stopped a clinic's learned attention transferring across the language its patients used. Still unbuild: translation, non-English summary generation, every language but Malay. The fourteen terms need native-speaker and clinical review — the mechanism is proven, the word list is a demonstration. Testing it also surfaced that **negation is unhandled in both languages** ("Patient denies chest pain" tags chest pain); pre-existing, now pinned in both so a fix cannot be applied to one and not the other.

Deferred: multilingual summary generation and handwriting capture (D-019). Summary generation in a second language is a *time* deferral only. Handwriting OCR was deferred **structurally**: a different ingestion pipeline (`image` → `OCR` → `redact` → `summarise`) where redaction is materially harder on noisy output — one mis-recognised character defeats a pattern that would have caught an identifier — and medical handwriting OCR is hard even for well-resourced products. Ambient voice capture serves the same "fast unstructured capture" need more safely.

Defects found after the build was "done", and what they have in common. The final pass found every enum column is declared `Mapped[StrEnum]` but backed by `String`, so reloaded rows return plain `str` and three `is` comparisons were silently dead branches (D-055) — most visibly, superseded highlight suggestions were never deleted and the Top Card rendered every claim twice. Four more surfaced afterwards from someone simply using the thing (D-059–D-062): a clinician's hand-marked span vanished from the Glance View; confirming one suggestion made every other suggestion on the open card return 404; "new since your last visit" stayed empty for an entire first session and then stopped advancing for anyone who refreshed often; and a task could be raised and never closed. A fifth, found while reproducing those: every timestamp left the API without a UTC offset, so a browser read it as local time and a note written seconds ago rendered as "8h ago" in the timezone this was demoed in.

All of them survived a green suite, and for the same reason. Each lives in the seam between two pieces of individually correct code. The manual-highlight bonus is right where it is written and right where it is recomputed; only the ordering of the two is wrong. The timestamps are right in the database and right in the browser; only the contract between them was unstated. Component-level tests cannot see this class, which is why the regressions are written as end-to-end sequences — open the chart, write a note, reload — and why ten of the fifteen fail against the previous commit. A test that has never seen its regression is a description of current behaviour wearing a test's clothing.

Reported at this length because the honest version of "structural enforcement catches mistakes" has to include the classes it missed. The enum column-type migration and a real end-to-end browser test are both deferred rather than attempted hours before submission.

7. Security posture

One table, three statuses, nothing hedged. It is longer than a list of things that work would be, because about a third of the rows are gaps rather than features. We would rather a reviewer see the gaps here than find one we had not mentioned, since an undisclosed gap suggests there are others. Everything marked "known gap" is something that would need fixing before this went near a real patient, and the paragraph below the table says plainly why it is not ready for one.

Implemented = built, tested, verifiable here. **Documented decision** = deliberate for a 72-hour prototype, production shape stated. **Known gap** = genuinely missing, listed because a control whose edges nobody knows is worse than a weaker one everybody understands.

Area	Status
PHI redaction chokepoint	Implemented — regex + gazetteer, fail-closed
Stored-XSS / content safety	Implemented — never rendered as HTML, enforced by source scan
RBAC (role + clinic, server-side)	Implemented — fused, proven over HTTP
Logging hygiene	Implemented — content-free by construction, verified by grep
JWT storage	Implemented — <code>httpOnly</code> cookie, <code>SameSite=lax</code> , 60min TTL
AI note immutability; comment isolation from patients	Implemented — corrections supersede; comments refused at route <i>and</i> stamped internal
Conflict handling	Implemented — precedence <i>and</i> flag; disputed content never deleted
Audio never persisted; un-redacted egress	Implemented — memory only, asserted against the DB; egress gate fails closed, never degrades to stub
Learning substrate holds no prose; clinic-partitioned; safety floored	Implemented
Enum comparison correctness	Implemented (guarded) — <code>==</code> throughout; source scan fails the build (D-055)
CSRF defence	Documented decision — <code>SameSite=lax</code> only
TLS in transit	Documented decision — terminates at reverse proxy / LB with HSTS in production; not implemented locally (plain HTTP)
Encryption at rest	Documented decision — managed-Postgres volume encryption or SQLCipher in production; SQLite here is unencrypted
Password hashing; decay scheduling	Documented decision — PBKDF2 120k (argon2id in prod); explicit trigger, no cron
Token refresh / rotation / revocation	Known gap — no refresh flow, no denylist
Login rate limiting	Known gap
Redaction recall on unanticipated names	Known gap — lowercase/transliterated names in prose can survive
Scribe failure recovery	Known gap — synchronous; a crash mid-run loses the summary
Real speech recognition	Known gap — default recogniser is a simulated stub ; no audio has ever been transcribed by this build
Acoustic diarisation; consent artefact on patient recordings	Known gap — labels come from the transcript source; the clinician is a party and is never asked
Per-user normalisation of learning signals	Known gap — one enthusiast counts as consensus
Enum columns typed <code>String</code> not <code>Enum</code>	Known gap — structural fix for D-055
Formal accessibility / WCAG audit	Known gap — colour is never the sole signal, but no audit was run
Risk ordinal floored by deterministic rules	Implemented — a model may raise a level, never lower one; provenance stored per note (D-066)
Patient-facing generation	Implemented — structurally impossible, guarded at import, not gated by an approval step someone can click through (D-067)
Human-human contradiction detection	Implemented (narrow) — allergy/dose/status, deterministic, never auto-resolved (D-068)
Contradiction recall	Known gap — watchlist not formulaic; absence of a flag is not evidence of agreement
Confidence calibration	Known gap — hedging density is a proxy, never validated against labelled data
Exposure bias in the learning loop	Partially mitigated — one reserved exploration slot; no off-policy evaluation (D-069)
Wire-format timezone correctness	Implemented (by convention) — UTC offsets via one annotation, pinned by tests that walk payloads; a new endpoint can still regress it (D-061)
End-to-end browser testing	Known gap — component tests mock <code>Api</code> , so a change to the fetch layer is caught by neither suite

Plainly: locally there is no TLS and no encryption at rest — plain HTTP on localhost, an unencrypted gitignored SQLite file. Both are deployment configuration rather than application code, hence decisions with a stated production shape; the honest consequence is that **this build is not safe for real PHI as-is**, which the README states too so it cannot be missed by someone who opens one file.

Verification. 486 backend tests plus 25 frontend component tests, no API key or network needed. Access-control and history tests were **deliberately broken to confirm they can fail** — reversing D-004 fails exactly the staff-visibility tests, removing the clinic filter fails 15, disabling the conflict guard fails 3. The same discipline applied to the Phase 7 regressions: 10 of 25 frontend tests and 10 of 15 backend ones fail when the code they cover is reverted. A security test that cannot fail is worse than none, because it is mistaken for coverage. Logs were grepped for planted names, identifiers and body text after exercising every route: zero hits **on the success path**. The failure path was a different story, and the reviewers' scenario 3 is what found it: an unhandled exception reached Starlette's default handler and unicorn logged the traceback, and SQLAlchemy embeds bound parameters in its exception messages — so a name, an NRIC and note content landed in one `stderr` line. `log_event` made it hard to log content *on purpose*; nothing covered content logged on our behalf by code we did not write. Fixed in D-071 and pinned by three tests that fail without the middleware. The original claim was true and the wrong question had been asked of it.

Round two of the review — the sixteen clinic scenarios, the twelve-capability list, what we tried that failed, and which earlier assumptions no longer stand — is a separate document: [BRIEF_ROUND2.md](#).